

# Deep RL Arm Manipulation

The goal to the Deep RL Arm Manipulation project is to create a DQN agent and define reward functions to teach a robotic arm to carry out two primary objectives:

1. *Have any part of the robot arm touch the a designated object on the ground, with at least a 90% accuracy.*
2. *Have only the gripper base of the robot arm touch the designated object, with at least a 80% accuracy*

The template project is based on the Nvidia open source project "jetson-reinforcement" developed by Dustin Franklin.

## Introduction

Deep Reinforcement Learning is an active area of research in robotics. This project is a challenging Deep RL implementation using a gazebo environment and an arm plugin (ArmPlugin.cpp) that assigns rewards, either **positive** or **negative**, based on simple rules. During repeated attempts by the robot to touch a cylinder in the gazebo world without touching the ground, the goal is to train the robot arm to reach the desired accuracy goal. The project was implemented on a Jetson TX2 board running Linux 16.04. The project sources reside in the git directory : <https://github.com/mbufi/RoboND-DeepRL-Project>

## Reward functions

The reward functions are defined in [ArmPlugin.cpp](#). The arm joints were updated using position control (default settings) instead of using velocity control. For each of the joints, there are two actions; increase or decrease the joint position.

- **REWARD\_WIN** was set to **0.125** ( **0.1** for the second objective) with **REWARD\_LOSS** to **-0.125** ( **-0.1** for the second objective).
- If the robot gripper hit the ground, a **REWARD\_LOSS \* 10** was given and the episode ended.

Within the episode, interim rewards were issued if there was no ground contact or 100 frames had not been exceeded.

The main interim reward of each episode was based on the distance goal delta between the gripper and the Cylinder prop. If a positive weighted average was derived then a **REWARD\_WIN** was recorded otherwise **REWARD\_LOSS \* distance\_to\_goal** was issued. Thus the **REWARD\_LOSS** would be higher the further the arm was from the goal.

An additional **REWARD\_LOSS** was added in the second objective for the gripper base stating:

- if the absolute average goal delta was **< 0.001** to penalize no movement.

For the first objective:

- If the robot arm hit the prop, a **REWARD\_WIN \* 10** was used

For the second objective:

- **REWARD\_LOSS \* 5** if the collision was not with the **gripper\_middle**.
- **REWARD\_WIN \* 20** was issued if the collision point was **gripper\_middle**.

For both objectives , **any collision ends the episode**.

## **Hyper Parameters**

Image dimensions were set to the same size as the input. Training was performed on a the Jetson TX2, which required no fine tuning to improve performance. INPUT\_WIDTH 64 x INPUT\_HEIGHT 64

The OPTIMIZER was changed from none to Adam. Possible common options are Adam or RMSProp, which are instances of adaptive learning rates to use in gradient-based optimization of parameters. Adam slightly outperforms RMSprop, which in turn gave good results. The parameter use LSTM was set to true to enable the LSTM agent

For objective 1 the LEARNING\_RATE was 0.1 with REPLAY\_MEMORY at 1000 . The value was chosen via trail and error.

For objective 2 the LEARNING\_RATE was decreased to 0.01 due to the higher REPLAY\_MEMORY set at 20000 . The higher REPLAY\_MEMORY was used for more discrete learning, due to the smaller surface area required to achieve a collision to meet objectives.


BATCH\_SIZE was set to 512 (again sufficient memory on the TX2).

LSTM (Long short-term memory) was used USE\_LSTM = true with a size of LSTM\_SIZE = 256 which was set via trial and error.

In general, most parameters were set based on trail and error since I was using a TX2 and not a regular graphics card. TX2 is an experimental hardware platform, therefore optimization of parameters is extremely crucial for proper training.

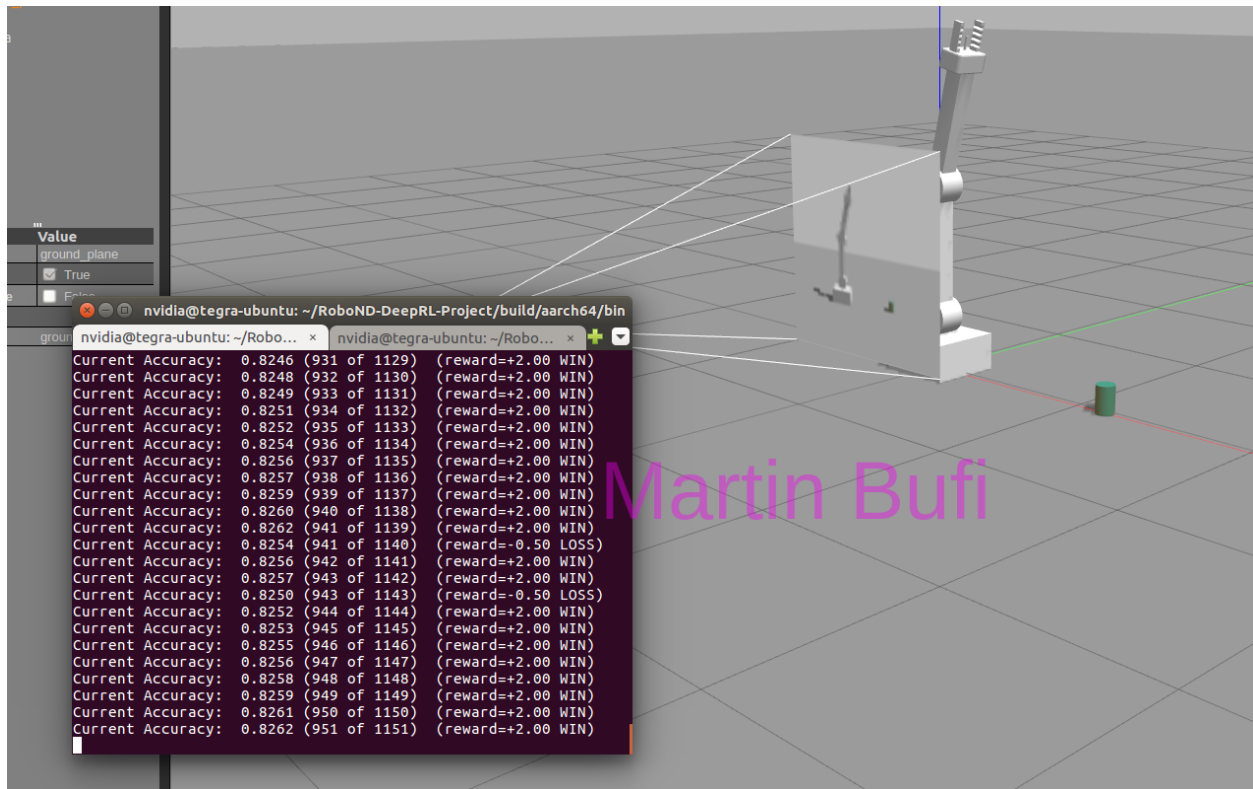
## **Results**

**Objective 1** : Have any part of the robot arm touch the object of interest, with at least a 90% accuracy for a minimum of 100 runs.



For reasons unknown, the robotic arm struggled at first to hit the prop. After some trial and error, it began learning how to hit it properly and with a degree of accuracy in a repeatable fashion. On occasion the arm would try a different path and begin losing its training. But due to having memory, it would then go back to using the winning sequence to hit the prop. As shown above, the objective was achieved well within the criteria specified with an accuracy of 92% before I stopped the script. I am sure it would have gone higher, if left to run longer.

**Objective 2:** Have only the gripper base of the robot arm touch the object, with at least a 80% accuracy for a minimum of 100 runs.



With the finer control required (only hitting the gripper based), and alteration to the interim reward system, this configuration would often hesitate before making a move. During the early stages of training, it learned how to get very close to the prop but never make it there since it was so hesitant. Once it learned how to effectively hit the prop with the middle of the gripper, it began experimenting other ways to do so. Such as extending past the prop and then coming back for it in an arch (as if it was digging), but this failed due to it usually hitting the prop with its arm joint. This whole process of learning the proper way to do it without hesitation took the arm over 100 iterations before becoming more confident. Decreasing the learning rate to 0.01 and increasing LSTM memory to 20000, were what allowed it to learn the precise movements to reliably complete the objective.

## Discussion

Training a DeepRL agent is a time consuming task, often taking many hours of computer time. Tuning parameters seems like black magic. The variability from run to run seems to far exceed the effect of any parameter changes. Tuning parameters such as LSTM size can give different results, some of them amusing and unexpected. Perhaps using a different approach to the problem, such as using a genetic algorithm would allow for less manual labor of tuning hyper parameters.

## Conclusion and Future Work

This C++ API was developed by Nvidia and optimized specifically for embedded platforms such as the TX2 performed extremely well on the native TX2 platform. The code took advantage of the Jetson CUDA architecture and ran the episodes at a rate of about one per second. Building the project on the Jetson platform was very simple and it performed well. With some fine tuning, one would not notice any performance degradation in such a small platform.

Due to the randomness of the algorithm in this project, one begins to wonder if this randomness could be reduced or channeled by looking at Genetic Algorithms for Deep RL. The idea is that populations of living organism (or in this case, the arm) learn to adapt to their environment by a means of natural selection. Perhaps something such as this could be used to give the DeepRL agent more guidance in choosing parameters to accomplish the goals. This would allow the user (myself) to start up the script and allow the generations of arms to learn how to adjust the hyper parameters themselves, leaving the developer to sit back and watch the magic unfold. This is something I would like to further research and perhaps implement on the TX2.